

Great—you're right at the point where most people either get it working... or get stuck. Let's walk it cleanly from your **edited Makefile** → **running TPC-H on MySQL**.

Step-by-step from where you are

1. Save and exit the Makefile

If you're in nano:

- Press CTRL + O → Enter (save)
 - Press CTRL + X (exit)
-

2. Compile dbgen

Still inside dbgen/, run:

make clean

make

3. Verify build worked

Run:

ls

You should see:

dbgen

qgen

If not, stop here and tell me the error.

4. Generate data (start small)

./dbgen -s 1

This creates files like:

customer.tbl

orders.tbl

lineitem.tbl

...

5. Fix the .tbl format (IMPORTANT for MySQL)

TPC-H adds a trailing | at the end of each row—MySQL hates that.

Run this:

```
for f in *.tbl; do sed 's/|$/ /' "$f" > "clean_$f"; done
```

Now you'll have:

clean_customer.tbl

clean_orders.tbl

...

6. Open MySQL

```
mysql -u root -p
```

7. Create database

```
CREATE DATABASE tpch;
```

```
USE tpch;
```

8. Create tables (you need this)

TPC-H's default schema isn't MySQL-ready, so use this working version.

Example (run this first):

```
CREATE TABLE region (
```

```
  r_regionkey INT PRIMARY KEY,
```

```
  r_name CHAR(25),
```

```
    r_comment VARCHAR(152)
);

CREATE TABLE nation (
    n_nationkey INT PRIMARY KEY,
    n_name CHAR(25),
    n_regionkey INT,
    n_comment VARCHAR(152)
);

CREATE TABLE customer (
    c_custkey INT PRIMARY KEY,
    c_name VARCHAR(25),
    c_address VARCHAR(40),
    c_nationkey INT,
    c_phone CHAR(15),
    c_acctbal DECIMAL(15,2),
    c_mktsegment CHAR(10),
    c_comment VARCHAR(117)
);
```

👉 You'll need all 8 tables.

If you want, I can paste the full schema in one block.

9. Enable file loading

Inside MySQL:

```
SET GLOBAL local_infile = 1;
```

Exit and reconnect:

```
mysql --local-infile=1 -u root -p tpch
```

10. Load data

Example:

```
LOAD DATA LOCAL INFILE '/full/path/dbgen/clean_customer.tbl'
```

```
INTO TABLE customer
```

```
FIELDS TERMINATED BY '|'
```

```
LINES TERMINATED BY '\n';
```

Repeat for all tables:

- region
- nation
- supplier
- customer
- part
- partsupp
- orders
- lineitem

11. Quick test

```
SELECT COUNT(*) FROM lineitem;
```

If you see a big number →  success

12. Generate a query

Back in bash:

```
./qgen 1 > q1.sql
```

13. Run query in MySQL

```
mysql -u root -p tpch < q1.sql
```

Expected issues (don't panic)

With MySQL:

- Some queries may fail due to:
 - INTERVAL syntax differences
 - date arithmetic
 - You may need small edits
-

Where you are now

You just finished the hardest part:

✓ Makefile edited

✓ dbgen compiled

☐ data generated

☐ data loaded into MySQL

☐ queries running

If you want next

I can:

- give you **all 8 MySQL tables (copy-paste ready)**
- fix **TPC-H queries for MySQL compatibility**
- or show a **much easier route using DuckDB (1 command, no schema pain)**

Just tell me 